

BCSE498J- Project - II

DETECTION AND MITIGATION OF JAILBREAK ATTACKS IN LLM S

Submitted in partial fulfilment of the requirements for the degree of

Bachelor of Technology

in

Computer Science and Engineering

by

22BCE3689

DIVYESH NARSHI KANZARIYA

Under the Supervision of

Prof. Umamaheswari M

Assistant Professor Sr. Grade 1

School of Computer Science and Engineering




VIT[®]
Vellore Institute of Technology
(Deemed to be University under section 3 of UGC Act, 1956)

April, 2026

DECLARATION

I hereby declare that the project entitled "DETECTION AND MITIGATION OF JAILBREAK ATTACKS IN LLM S" submitted by me, for the award of the degree of Bachelor of Technology in Computer Science and Engineering to VIT is a record of bonafide work carried out by me under the supervision of Prof. Umamaheswari M, School of Computer Science and Engineering.

I further declare that the work reported in this project report has not been submitted and will not be submitted, either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university.



Place: Vellore

DIVYESH NARSHI KANZARIYA (22BCE3689)

Date : 29/04/2026


CERTIFICATE


This is to certify that the project entitled "DETECTION AND MITIGATION OF JAILBREAK ATTACKS IN LLM S" submitted by DIVYESH NARSHI KANZARIYA (22BCE3689), School of Computer Science and Engineering, VIT, for the award of the degree of Bachelor of Technology in Computer Science and Engineering, is a record of bonafide work carried out by the student under my supervision during Winter Semester 2025-2026, as per the VIT code of academic and research ethics.


The contents of this report have not been submitted and will not be submitted either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university. The project fulfills the requirements and regulations of the University and in my opinion meets the necessary standards for submission.

Place: Vellore

Date : 28/04/2026


Internal Examiner


Signature of the Guide


External Examiner



ACKNOWLEDGEMENT

I DIVYESH NARSHI KANZARIYA (22BCE3689) am deeply grateful to the management of Vellore Institute of Technology (VIT) for providing me with the opportunity and resources to undertake this project. Their commitment to fostering a conducive learning environment has been instrumental in my academic journey. The support and infrastructure provided by VIT have enabled me to explore and develop my ideas to their fullest potential.

My sincere thanks to Dr. Jaisankar N, the Dean of the School of Computer Science and Engineering, for constant support and encouragement. The Dean's leadership and vision have greatly inspired me to strive for excellence. The dedication to academic excellence and innovation has been a constant source of motivation.

I express my profound appreciation to Dr. Boominathan P, the Head of the Department of Software Systems for insightful guidance and continuous support. The expertise and advice have been crucial in shaping the direction of my project. The commitment to fostering a collaborative and supportive atmosphere has greatly enhanced my learning experience. The constructive feedback and encouragement have been invaluable in overcoming challenges and achieving my project goals.

I am immensely thankful to my project supervisor, Prof. Umamaheswari M, School of Computer Science and Engineering, for dedicated mentorship and invaluable feedback. The patience, knowledge, and encouragement have been pivotal in the successful completion of this project. The willingness to share expertise and provide thoughtful guidance has been instrumental in refining my ideas and methodologies. This support has not only contributed to the success of this project but has also enriched my overall academic experience.

I extend my heartfelt thanks to all for the guidance and support received throughout this endeavor.

TABLE OF CONTENTS

Executive Summary	ii
List of Tables	iii
List of Figures	iv
List of Abbreviations	v
List of Symbols	vi
1 INTRODUCTION	1
1.1 Background	1
1.2 Motivation	1
1.3 Scope of the project	2
2 PROJECT DESCRIPTION AND GOALS	5
2.1 Literature review	5
2.2 Research Gap	6
2.3 Objectives	7
2.4 Problem statement	8
2.5 Project plan	8
3 TECHNICAL SPECIFICATION	12
3.1 Requirements	12
3.1.1 Functional	12
3.1.2 Non-Functional	12
3.2 Feasibility Study	13
3.2.1 Technical Feasibility	13
3.2.2 Economic Feasibility	13
3.2.3 Social Feasibility	14
3.3 System Specification	14
3.3.1 Hardware Specification	14
3.3.2 Software Specification	14

4	DESIGN APPROACH AND DETAILS	16
4.1	System Architecture	16
4.2	Detailed Design.....	17
4.2.1	Data Flow Diagram	17
4.2.2	Class Diagram	17
5	METHODOLOGY AND TESTING	19
5.1	Problem Definition and Threat Boundary Finalization.....	19
5.2	System Architecture and Defense Pipeline Architecture.....	19
5.3	Mathematical Formulation of Instructionality and Trust.....	20
5.4	Dataset Construction and Experimental Setup	20
6	PROJECT IMPLEMENTATION	22
6.1	Instruction-Data Boundary Detection Module	22
6.2	Graph-Based Cross-Chunk Collusion Analysis.....	22
6.3	Context Trust Scoring and Weighted Injection.....	23
6.4	Database & Vector Store Implementation	23
6.5	Software Integration and Frameworks.....	24
6.6	Five-Layer Defense Pipeline Implementation	24
6.6.1	Layer 1: Ingestion Sanitization and Provenance Verification.....	24
6.6.2	Layer 2: Access Controlled Search	24
6.6.3	Layer 3: Strengthening Prompt Structure (Spotlighting).....	24
6.6.4	Layer 4: Embedding Anomaly Detection.....	25
6.6.5	Layer 5: Output Validation and Monitoring	25
7	RESULTS AND DISCUSSION	26
7.1	Experimental Overview.....	26
7.2	Main Benchmark Results	26
7.3	Threshold Calibration and Score Separation	27
7.4	Latency and Runtime Overhead	29
7.5	Ablation Study	29
7.6	Stability Across Random Seeds	30
7.7	Adversarial Family Analysis	31
7.8	Graph-Based Collusion Analysis	32
7.9	Discussion	32

8	CONCLUSION AND FURTHER ENHANCEMENTS	34
8.1	Conclusion.....	34
8.2	Limitations	34
8.3	Further Enhancements.....	35

EXECUTIVE SUMMARY

This research proposes a modular, model-independent runtime security system that aims to increase RAG safety by mitigating indirect prompt injection, cross-chunk collusion, and jailbreak attacks on RAGs. Due to their dependence on outside information resources, language models have become exposed to maliciously inserted information during document retrieval. In this regard, security measures like prompt filtering and fine-tuning are frequently ineffective in such scenarios.

In order to overcome these shortcomings, the model proposes a five-step defense mechanism that works entirely during inference. The approach takes a multi-step process, involving instruction identification, graph-based collusion, trust-based contextual injection, and output groundedness. This ensures detection and mitigation of any potential threats through the process of content retrieval. The content chunk is analyzed to detect any instruction and relationships within. This allows the model to detect attacks, whether explicit or distributed. Once done, the system assigns a trust score to control the influence of each chunk while constructing the prompt.

Performance analysis is carried out using an experimentally designed set of data comprising good, poisoned, and colluding files, as well as adversarial queries. The findings show that there is clear differentiation between benign and malicious files, and the model performs consistently with fewer false alarms whenever the threshold is optimized. However, there are still weaknesses in dealing with paraphrasing and obfuscation techniques.

In general, the project proves that runtime protection for RAG systems is viable and functional in a specific context. It presents an innovative solution that can be extended to accommodate other needs, and it also highlights areas that need more attention in the future.

LIST OF TABLES

3.1	Hardware Specification	14
3.2	Software Specification.....	15
5.1	Dataset Mapping	21
6.1	Trust-Based Injection Strategies.....	23
7.1	Performance Comparison of Defense Mechanisms.....	27
7.2	Tuned Validation Operating Point Metrics	28
7.3	Latency Performance Metrics.....	29
7.4	Ablation Study of Defense Layer Contributions	30
7.5	Multi-Seed Stability Evaluation	31
7.6	Per-Category Adversarial Family Results	31

LIST OF FIGURES

2.1	Gantt Chart	11
4.1	System Architecture	16
4.2	Data Flow Diagram	17
4.3	Class Diagram	18
7.1	Threshold-Sensitivity Curve for Instructionality Threshold Selection	27
7.2	Score Calibration Plot	28
7.3	Calibrated ROC and Precision-Recall Curves	28
7.4	Latency Distribution and Runtime Overhead	29
7.5	Ablation Study Across Defense Layers.....	30
7.6	Interaction Graph Visualization.....	32

LIST OF ABBREVIATIONS

ACE	Attacks using Custom Encryptions
AI	Artificial Intelligence
API	Application Programming Interface
AUC	Area Under the Curve
BERT	Bidirectional Encoder Representations from Transformers
CoD	Chain-of-Detection
ControlNet	Firewall for RAG-based LLM Systems
DAN	Do Anything Now
DFD	Data Flow Diagram
FAISS	Facebook AI Similarity Search
GPT	Generative Pre-trained Transformer
HarmBench	Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal
HILL	Hypotheticality-based jailbreak attacks on Large Language Models
LACE	Layered Cipher jailbreak Attacks
LLM	Large Language Model
ML	Machine Learning
NLP	Natural Language Processing
PAIR	Prompt Automatic Iterative Refinement
PR	Precision-Recall
PromptBench	Benchmark for Evaluating Jailbreak Robustness
RAD	Retrieval-Augmented Defense
RAG	Retrieval-Augmented Generation
RLHF	Reinforcement Learning from Human Feedback
ROC	Receiver Operating Characteristic
TAP	Tree of Attacks with Pruning
TTV	Time-to-Violation

SYMBOLS AND NOTATIONS

$C(c_i)$	Collusion risk score of the i -th retrieved chunk
c_i	The i -th retrieved document chunk
$E(c_i)$	Embedding vector of the i -th chunk
$I(c_i)$	Instructionality score of the i -th retrieved chunk
$I_{centroids}$	Set of instructional centroid vectors
$L(c_i)$	Linguistic instructionality signal
W_{ji}	Semantic edge weight between chunks j and i
α	Weight assigned to semantic relevance in the instructionality score
β	Weight assigned to pedagogical relevance in the instructionality score
$\text{sim}(\cdot, \cdot)$	Cosine similarity function
$G = (V, E)$	Cross-chunk interaction graph
w_1	Weight assigned to answer consistency in trust score computation
w_2	Weight assigned to pedagogical confidence in trust score computation
w_3	Weight assigned to citation or source reliability in trust score computation

CHAPTER 1

INTRODUCTION

1.1 BACKGROUND

There have been fast-paced developments in the area of Artificial Intelligence (AI) through the development of Large Language Model (LLM) that understand, generate, and reason over natural language. Some of these models include Generative Pre-trained Transformer (GPT), Bidirectional Encoder Representations from Transformers (BERT) and Llama which have been used in several applications such as conversation agents, automated text analysis, question answering, and decision-making. These models involve pre-training over huge amounts of textual data to make them proficient enough in capturing linguistic phenomena and semantic relations.

Rule-based approaches for Natural Language Processing (NLP) were traditionally followed using manually constructed grammars and knowledge bases specific to particular domains. Deep learning with the use of transformers revolutionized the field where the model can be trained from the data without any hand-tuned parameters. Nonetheless, even with the superior performance levels of LLM, issues like hallucinations and inaccurate facts due to parameterized knowledge had been noted in the process of learning from data. The introduction of the approach called Retrieval-Augmented Generation (RAG) became important to address this problem

Nevertheless, although RAG has introduced various benefits, including knowledge-grounding capabilities and adaptability, several security threats have been identified through recent research in this architecture. Techniques such as prompt filtering, content moderation, and fine-tuning may not be effective in dealing with the threats of prompt injection and zero-day jailbreaking. Moreover, solutions that involve looking into the model's internal operations or training a guard model cannot be employed on API-based LLMs with limited source codes.

1.2 MOTIVATION

Integration of LLM into many key industries such as intelligent assistants, research platforms, healthcare systems, and enterprise systems has raised a lot of debate regard-

ing the security problems associated with these technologies. Although the usage of RAG improves the groundedness of the models, recent studies show that they also exhibit many vulnerabilities. This is because hidden instructions contained in the retrieved documents guide the model in behaving in a particular manner, without requiring any explicit user inputs.

Regarding the scope of this study, from a disciplinary standpoint, it fits into the field of artificial intelligence, Machine Learning (ML), cybersecurity, and natural language processing. It reveals the security risk created by the interaction of the three processes involved – retrieving the documents, context generation, and modeling. Regarding my interest in this project, it was inspired by my desire to leverage safety protocols of AI technology.

Based on the latest trends, there is an increased preference toward using black box Application Programming Interface (API) LLM, where the inner details of the model and its parameters are not available. This would make it impossible to perform any tuning as well as the inability to exploit the use of any model-specific defenses. Moreover, owing to an increase in zero-day jailbreak attacks and prompt injection, static filters will no longer be useful in tackling these types of attacks. The motive of this research study is to design a runtime defense mechanism that enhances security without having to change the structure of the model.

1.3 SCOPE OF THE PROJECT

1. Functional Scope: The functional scope of this project involves creating a multi-stage runtime defense pipeline that focuses on protecting LLM-based RAGs from manipulation by detecting adversarial instructions in the retrieved documents. Each retrieved document chunk is analyzed and categorized as factual information or hidden instruction. Based on the results, the instructionality score is calculated for each document. Additionally, interactions between multiple documents are analyzed semantically and logically to discover the possibility of instruction chains between documents and determine the risk of collusion between them.

Instead of utilizing binary filters based on allowing or blocking, the system assigns a trust score to each document and controls the extent to which each piece of information affects the LLM based on its instructionality score, collusion risk, and provenance

metadata. The whole process takes place externally to the language model without accessing its weights and logits, making it compatible with API-based models. Lastly, the generated output is checked for its groundedness, and when violations occur, the generation will be repeated under strict control or prevented entirely.

2. Non-Functional Scope: The model agnostic design of the system allows its seamless integration within various open-weight and API-based LLM. The system ensures that it has minimal latency overhead appropriate for interactive applications by restricting the number of hops in the cross-chunk interactions, making sure that the complexity does not increase exponentially with an increasing number of chunks. The usability of the system is assured as it guarantees that the number of false positives in benign queries is kept low, so that the defense mechanism does not interfere with the functioning of the rest of the system. The system is modifiable in nature, enabling replacement or enhancement of any of its individual parts such as the boundary detector, collusion detector, and trust scorer.

3. Technical Scope: Technical scope is restricted to text-based RAG pipelines only. Other inputs like images and audio signals are not covered under this project. The technology utilized in the system includes vector representations, semantic similarity calculations, interaction graph, and rules-based feature detection techniques. The system performance is evaluated based on the use of controlled datasets, artificial poison files, and sample queries that aim to test prompt injection and cross-chunk collusion attacks. Quantitative measures for assessing effectiveness include jailbreak success ratio, false positives rate, time taken until violation of policies, and groundedness deviation. The implementation is developed using Python 3.10 programming language, LlamaIndex as the framework for building RAG systems, Facebook AI Similarity Search (FAISS) for vector store, Sentence-Transformers for embeddings, NetworkX for graphs, and HuggingFace Transformers for LLM engine.

4. Project Boundaries:

Assumptions

- The defense layer functions completely independently from the LLM; no access to its inner workings, parameter values, or logits is presumed or needed.
- Testing involves artificially created poisoned documents and controlled benchmark searches, emulating potential attack vectors in practice.
- The design presumes that retrieved documents can serve as adversarial inputs, considering instructional risk in an accumulative sense rather than locally.
- The architecture is considered expandable for further development of features including multimodal defense systems, safeguarding long-term conversation history, and automated red teaming support.

Constraints

- The framework **does not involve** fine-tuning or training the language model.
- **Deployment at scale for production purposes** and optimizations for real-world use cases are not considered.
- The evaluation **is limited to** text-based retrieval-augmented generation pipelines; other modalities such as images or audio are excluded.
- Analysis of interactions across chunks is restricted in scope to avoid excessive runtime complexity.
- The defense mechanism **cannot prevent** an adversary from fine-tuning the model or poisoning its training dataset.

CHAPTER 2

PROJECT DESCRIPTION AND GOALS

2.1 LITERATURE REVIEW

LLM security has increasingly shifted from static alignment toward adversarial runtime defense, requiring methods that can detect prompt injection, scenario drift, and cross chunk collusion. Early work has identified Reinforcement Learning from Human Feedback (RLHF) Ouyang (2022) and red teaming Ganguli (2023) as the main strategies for aligning large language models with human values. Nevertheless, it has been noted that such safety mechanisms are surface level techniques, commonly referred to as being "a few tokens deep" Qi et al. (2024). This limitation is further evidenced by adversarial prompting techniques capable of circumventing aligned language model safeguards Carlini (2024). In the case where such security barriers are breached, safety recovery is needed to bring the LLM back to safe conditions Liu (2024b). There have been comprehensive surveys covering all aspects of LLM jailbreaking Xu et al. (2024), Yi et al. (2024).

From manual Do Anything Now (DAN) prompts for adversarial jailbreaking, Shen et al. (2024a), adversarial jailbreaking has evolved to automated large scale promotion of jailbreaks. Notable technologies used include the deployment of universal and transferable adversarial suffixes by Zou et al. (2023) and Zou (2025) and automation platforms such as the AutoDAN framework proposed by Liu et al. (2023), Prompt Automatic Iterative Refinement (PAIR) algorithm designed by Zhang et al. (2024), and Tree of Attacks with Pruning (TAP) method proposed by Mehrotra et al. (2023). The fast adaptive prompt technique has been discussed by Paulus et al. (2024), while Xiao (2024) proposed the use

Advanced approaches leverage obfuscation and vulnerabilities in structural models. Examples of these are attacks based on hypotheticality (Hypotheticality-based jailbreak attacks on Large Language Models (HILL)) Backes et al. (2024), custom encryption (Attacks using Custom Encryptions (ACE)) Shen et al. (2024b), stacked encryption schemes (Layered Cipher jailbreak Attacks (LACE)) Shen et al. (2025), and the exploitation of special tokens that are unique to a particular model Henderson et al.

(2024). Threats related to security are present in multi turn interactions via scenario drift Chen et al. (2024) and multilingual settings Yin (2025). Recent research efforts have expanded the threat landscape for multimodal LLM Li (2025b). Besides direct user input, prompt injection attacks Guha (2024) have been used to exploit the application of the model Agarwal (2024) and tool assisted agents Klyman (2024)

The emergence of Retrieval Augmented Generation (RAG) and AI agents brings about new types of vulnerabilities. RAGs may encounter issues such as cross chunk collusion attacks Li (2025a), groundedness biases Deng (2024), and trust weighted context injection Wang (2025b). Defensive measures against the former architecture include context isolation Chen et al. (2025), instruction–data ambiguity Shi et al. (2025), and specialized firewalls, such as ControlNet Yao et al. (2025). As for agents, their security can be enhanced through least privilege principles Phan (2025), framework protections Ramakrishnan and Balaji (2025), and JBSHIELD Liu (2025).

The defensive techniques entail timely tracking and strong decoding. Examples are the Attention Tracker technique by Hung et al. (2025), Chain-of-Detection (CoD) technique by Wu and Zhang (2025), and hybrid multi agent detection by Qin et al. (2025). Architectural safeguards encompass behavioral firewalls by Backes (2025), Retrieval-Augmented Defense (RAD) by Rawat et al. (2025), behavioral anomaly detection by Wang (2025a), and jailbreak resistant by He (2025). Policy based guardrails by Shao (2025) and instructional authority models provide additional control, but they remain incomplete against adaptive jailbreak strategies and cross chunk attacks, while defense feedback loops Zhang (2025) tackle adaptive attacks on safety filters by Lyu et al. (2024). Lastly, the community benefits from benchmarking platforms such as HarmBench Mazeika et al. (2024), PromptBench Liu (2024a), LLM Fuzzer Yu et al. (2024), automated red teaming Fredrikson (2024), and black box jailbreak benchmarking Picek (2024)

2.2 RESEARCH GAP

Despite the extensive research conducted, several significant gaps remain in current LLM security frameworks:

- **Vulnerability of RAG Multi Chunks:** Although context isolation is available in Chen et al. (2025), present implementations cannot prevent cross chunk coor-

dination threats that distribute malicious instructions in more than one extracted document Li (2025b), resulting in groundedness anomalies Deng (2024).

- **Delay of Real Time Monitoring:** Modern protective techniques including hybrid multi agent detection Qin et al. (2025) and Chain of Detection Wu and Zhang (2025) tend to have inherent limitations in terms of speed, thereby making their deployment impractical for high traffic environments compared to retrieval enhanced security mechanisms Rawat et al. (2025).
- **Evading Behavioral Controls through Encryption:** Behavioral firewall approaches Backes (2025)) and anomaly detection systems Wang (2025a) are ineffective against attacks carried out using customized encryption methods (ACE) Shen et al. (2024b) or multiple layers of cipher text (LACE) Shen et al. (2025), which prevent token analysis from identifying any malicious intent.
- **Alignment Depth:** The issue that alignment for safety is only "a few tokens deep" Qi et al. (2024) is another crucial problem since models can be susceptible to scenario drift during multi turn jailbreak attempts Chen et al. (2024) and require repeated training to regain safety alignment Liu (2024b).
- **Ambiguity of Agent Authority:** The current state of agent security does not have a common capability for instructing authorities in Chen (2025), which allows implementing the principle of least privilege Phan (2025) to protect against attacks based on hijacking agents Ramakrishnan and Balaji (2025), Liu (2025).
- **Insufficiency of Existing Benchmarks for Multimodality and Multiple Languages:** Standard benchmarks Harmbench Mazeika et al. (2024) and Prompt Bench Liu (2024a) cannot measure the effectiveness of multimodal jailbreaking Li (2025b) and multilingual attacks Yin (2025) against text filters.
- **Scalability of Automated Testing:** While there are tools such as LLM Fuzzer by Yu et al. (2024) and automated red teaming by Fredrikson (2024), they do not seem to match the speed of production of universal adversarial suffixes by Zou (2025) and adaptive attacks against safety filters by Lyu et al. (2024).

2.3 OBJECTIVES

- Firstly, one of the objectives of the study is the creation of a defensive runtime environment for LLM and RAG systems which would help detect zero day jailbreak attacks using behavioral traits such as instruction drift, authority escalation, and document retrieval anomaly.
- Secondly, the objective of this project involves protecting RAG systems against indirect prompt injections. This includes developing a method for separating instructive inputs from facts presented in the document, and detecting cross document collusion for forming malicious instructions.
- Some of the other objectives of this research include employing trust based injection over allow/block injection schemes for reducing false positives while retaining usability of the system. Finally, a generalization of the defense algorithm over multiple large language models would also be developed.
- In order to assess the efficacy of the proposed mechanism, numerical parameters such as jailbreak success rate, Time-to-Violation, false positive rates, and groundedness deviation would be used. Additionally, an evaluation framework would also be developed for the purpose.

2.4 PROBLEM STATEMENT

According to the most recent studies, LLM based systems are susceptible to zero day jailbreaks, multi turn instruction drifts, and indirect prompt injections. With indirect prompt injections, the malicious commands would be embedded into the document rather than typed in by the user, and hence they cannot be detected via input sanitization.

Essentially, the problem lies in the fact that the vulnerabilities created by instructions in RAG systems have an emergent and accumulative nature. Different benign chunks could become malicious when combined in a single window. Current measures to counteract the risk have only been considered from the aspect of analyzing individual chunks or fine tuning the models.

2.5 PROJECT PLAN

The task will be conducted systematically by dividing it into phases that span several weeks, ensuring a structured approach from problem definition to the evaluation and documentation phases. This plan ensures that each phase is logically dependent on previous ones, proper reviews are conducted on time, and the runtime defense system is developed incrementally.

Phase 1: Problem Definition and Scope (Week 1) The first phase of this task entails defining the problems related to Large Language Models and RAG pipelines. It involves specifying the project scope, objectives, assumptions, and threat model boundaries. A major emphasis is put on constraining the proposed defense strategy to the runtime level, without altering or fine tuning the underlying LLM model. Deliverables: Problem Statement, Project Scope, Threat Model

Phase 2: Literature Review and Research Gap Analysis (Week 2) An extensive review of the literature is performed, highlighting any existing research work on jail-break attacks, indirect injection of prompts, RAG weaknesses, and existing defense methods. The shortcomings of any existing solutions are discussed, pinpointing important research gaps like the lack of analysis involving interactions among chunks and considering trust while handling contexts. Deliverables: Literature review document, Research gaps identified

Phase 3: System Architecture Design (Weeks 3–4) The system architecture design for the runtime defense layer is carried out. The overall pipeline for instructions and data boundaries detection, interactions between chunks analysis, injection of context based on trust levels, and groundedness assessment is designed. Deliverables: System architecture diagram, Defense pipeline design

Phase 4: Attack Development and Dataset Creation (Weeks 5–6) In this phase, evaluation datasets are created, consisting of benign knowledge documents and artificial poisoning documents that mimic indirect prompt injection and cross chunk collusion attacks. Experimentally reproducible attack scenarios are generated for testing. Outputs: Artificial poisoning document set, Benign data set, Documentation for attacks

Phase 5: Instruction/Data Boundary Identification (Week 7) The instruction/data boundary identification component is programmed to examine the selected chunks and calculate the degree of instructionality present within each chunk. Outputs: Instruction/data boundary identification component, Instructionality score for each chunk

Phase 6: Analysis of Interaction Between Chunks (Week 8) This phase centers on creating models for relationships between retrieved chunks. An interaction model for chunks will be created to identify emerging instruction sequences and collusion risks while keeping the runtime complexity within acceptable levels. Deliverables: Interaction engine for chunks, Collusion risks assessment

Phase 7: Trust Driven Context Injection (Week 9) The instruction scores, collusion risk values, and provenance information about chunks will be used for computing the trust scores of chunks. Then, the trust driven context injection approach will be implemented that controls the authority of chunks. Deliverables: Trust engine, Context injection manager

Phase 8: End to end Defense Integration (Weeks 10 12) Integration of all defense components takes place into a unified runtime pipeline which is connected to the RAG system and LLM inference pipeline. Behavior independence from model specifics is checked. Deliverables: End to end defense pipeline in the runtime environment, Configurable RAG defense pipeline

Phase 9: Output Groundedness Verification (Weeks 13 15) Verification mechanisms after generation are added to make sure that output responses do not rely on untrusted content and can be safely regenerated or rejected otherwise. Deliverables: Deviation detector for generated outputs, Output validation mechanisms

Phase 10: System Evaluation and Performance (Weeks 16 18) Experimental evaluation of the developed solution against both benign and previously unknown jailbreak attacks is conducted and corresponding metrics (e.g., false positive rate, jailbreak success rate, time to violation, etc.) are collected and analyzed. Deliverables: Experimental results, Comparison charts

Phase 11: Project Reporting and Submission (Weeks 19 20) This phase is aimed at producing final project documentation that would include methodology, experimental setup, limitations of approach, and future research directions, as well as producing the final report and reproducible codebase. Deliverables: Project report, Presentation slides, Codebase




Figure 2.1: Gantt Chart

CHAPTER 3

TECHNICAL SPECIFICATION

3.1 REQUIREMENTS

The designed system has to analyze each individual chunk, along with chunks in their entirety, in order to identify whether they include any factual information, covert commands, or a combination of malicious activities altogether. Moreover, the system has to calculate an instructionality measure for each chunk, which would help it assess the probability of its contents being intended for changing model behavior instead of conveying information. Finally, the model has to evaluate semantic and logical connections between the extracted fragments in order to uncover any collusion and malicious activity.

3.1.1 Functional

Below is a description of the required functional capabilities of the proposed system.

The system will assess each retrieved document segment to ascertain if the information consists of facts, commands, or a combination of both. The system will measure the instruction score of each segment by applying semantic and linguistic analysis methods. The system will examine the relationship between several retrieved segments to detect collusion among the segments and distributed malicious commands. The system will measure the collusion score, which reflects how much instructions reinforce one another across related segments. Furthermore, the system will assign a trust score to each retrieved segment using the instruction score, collusion threat, and source origin. The system will perform trust-weighted context injection, meaning that suspicious documents will have little impact on the language model's processing phase. Besides, the system will function as a runtime layer without model retraining and will verify groundedness using retrieved documents.

3.1.2 Non-Functional

The non-functional requirements below specify the system quality and performance properties.

For non-functional requirements, usability, performance, and maintainability are of paramount importance. It should be model agnostic and support integration with open-weight as well as API-based LLMs. In addition, it should cause little overhead in terms of latency to ensure the usability of the defense for interactive applications. False positives on benign queries should also be low. Furthermore, the complexity should not exceed bounds through depth restrictions on cross-chunk analysis interactions. Reproducibility should be ensured via use of open-source libraries as well as publicly available models, while individual modules should be easily upgradable/replaceable. Lastly, it should scale relatively well in case of increasing number of chunks retrieved.

3.2 FEASIBILITY STUDY

Feasibility assessment was carried out with regard to the implementation of the proposed solution from the technical, economic, and social viewpoints.

3.2.1 Technical Feasibility

From the technical point of view, it can be asserted that the proposed solution is completely feasible, as its functioning does not require any special knowledge about model internals and only implies text analysis techniques. Specifically, the suggested approach utilizes well-developed algorithms such as vector embeddings, semantic similarity evaluation, graph interactions analysis, and other standard linguistic approaches. As a result, all calculations are computationally affordable due to the bounded graph depth.

3.2.2 Economic Feasibility

The proposed project is economically viable because it leverages open-source libraries and frameworks. Licensing fees for the usage of vector databases, embedding algorithms, graph processing libraries, and large language model inference runtimes do not apply to this system. Moreover, the application can be deployed on standard hardware, avoiding the need for expensive graphics processors for training and fine-tuning of models. Thus, the system is affordable to deploy in educational environments and smaller companies.

3.2.3 Social Feasibility

The proposed system is socially beneficial because it enhances the security and ethical nature of AI-based applications. Jailbreaking and indirect prompt injection attacks can be avoided with the use of this approach, preventing malicious actors from exploiting large language models in critical areas, including medical care, education, and corporate decision making.

3.3 SYSTEM SPECIFICATION

This section outlines the hardware and software requirements needed to implement and evaluate the proposed system.

3.3.1 Hardware Specification

The minimum hardware configuration required to run the system is as follows.

Table 3.1: Hardware Specification

Component	Specification
Processor	Intel Core i5 / AMD Ryzen 5 or higher
Ram	Minimum 8 GB (16 GB recommended)
Storage	Minimum 50 GB free disk space
GPU	Optional (for faster embedding computation); not mandatory
Operating System	Windows 10 / Linux / macOS

3.3.2 Software Specification

The software stack required for implementation and evaluation is listed below.

Table 3.2: Software Specification

Component	Specification
Programming Language	Python 3.10
RAG Framework	LlamaIndex
Vector Database	FAISS or Chroma
Embedding Models	Sentence-Transformers
Graph Processing	NetworkX
LLM Runtime	HuggingFace Transformers / llama.cpp
Development Environment	VS Code / Jupyter Notebook
Operating System	Windows / Linux

CHAPTER 4

DESIGN APPROACH AND DETAILS

4.1 SYSTEM ARCHITECTURE

The design of the system is intended as a modulated defense pipeline running during runtime. This system is external to the retrieval system but comes right after it and before the LLM’s inferencing engine. The system filters documents before they are added to the model’s context through various stages aimed at detection and filtering of instructions. In addition, the system works on individual chunks and performs a cross-check on them together with calculating their trust score, and validating the outputs to ensure their groundedness. Figure 4.1 provides an illustration of the system architecture.




Figure 4.1: System Architecture

4.2 DETAILED DESIGN

The system design is depicted using Data Flow Diagrams and Class Diagrams, which show how data flows within the defense pipeline system and how the different elements of the system are designed to function with each other.

4.2.1 Data Flow Diagram

A depiction of how the chunks of the document obtained will flow through the defense process is presented using the Data Flow Diagram (DFD). This includes the intake stage, where the documents are processed and scored for instructionality, then through cross-chunk collusion and trust-weighted contextual insertion, and finally to the controlled LLM inference and verification stages.




Figure 4.2: Data Flow Diagram

4.2.2 Class Diagram

The Class Diagram is used to show how the structure of the modules within the system is designed. In this regard, it contains information about the important classes in the system, including the features, behaviors, as well as dependency among the classes.

As per the modular design of the pipeline, there is a separate class for each of the defense components such as boundary detector, collusion analysis, trust score computation, context injection, and groundedness validation.




Figure 4.3: Class Diagram

CHAPTER 5

METHODOLOGY AND TESTING

The methodology chosen for this study involves creating a runtime security middleware whose functionality is strictly confined to its operation in a "black-box" fashion without any interaction with the inner workings of the LLM. Such an approach becomes a requirement due to the necessity of deploying the system within real-world situations involving the use of closed-source APIs or local models, which prevents access to signals such as attention scores and logit distribution. The new methodology moves beyond signature detection and onto behavior detection by analyzing the structure and semantics of retrieved context.

5.1 PROBLEM DEFINITION AND THREAT BOUNDARY FINALIZATION

The first step of this methodology includes precise definition of the problem domain and finalization of threat boundaries. According to the research, there are two major threats that can be posed against the RAG pipeline. Firstly, the problem involves indirect prompt injection when just one document retrieved includes instructions for prompt overriding. Secondly, a more complicated threat comes in the form of cross-chunk collusion which is spread between several different documents. The assumptions for this method include a highly available environment in which the vector database might be compromised by being poisoned with the external material or including untrusted content obtained through web scraping or even uploading. The boundary of protection will be made on the inference-time interface of this pipeline, making it five stages.

5.2 SYSTEM ARCHITECTURE AND DEFENSE PIPELINE ARCHITECTURE

The architecture is built on a pipeline system to facilitate data flow at runtime. When a user initiates an interaction with the system, the RAG retriever retrieves relevant chunk information from the vector database, which is then processed through the defense layer architecture before formatting the information to the LLM. There are five key steps involved in the defense pipeline process. The first step is detecting the instruction-data boundaries; second, cross-chunk collusion detection; third, trust scoring; fourth,

context construction using trust weight; and fifth, groundedness deviation detection. Each of these steps is engineered using best practices that include the least privilege principle.

5.3 MATHEMATICAL FORMULATION OF INSTRUCTIONALITY AND TRUST

In order to create a deterministic and testable defense mechanism, a mathematically based system is required to measure the risk associated with each document. In other words, this involves the computation of an instruction score ($I(c_i)$) for every chunk c_i obtained in the search, which is the likelihood that the chunk tries to alter the behavior of the model rather than providing facts. The score is formulated using a weighted hybrid approach:

$$Instructionality(c_i) = \alpha \cdot \text{sim}(E(c_i), I_{\text{centroids}}) + \beta \cdot L(c_i) \quad (5.1)$$

Where, $E(c_i)$ indicates embedding vector for the i -th chunk. $I_{\text{centroids}}$ indicate set of instructional centroids obtained based on existing jailbreak and injection patterns. $\text{sim}(\cdot, \cdot)$ refers to cosine similarity. $L(c_i)$ indicates linguistic score signal computed on the basis of occurrence of imperative verbs and authority-shifting cues. This is supplemented by the Collusion Risk ($C(c_i)$), which is propagated through a directed graph $G = (V, E)$ where nodes are retrieved chunks and edges represent semantic reinforcement. The risk for a node i is calculated as:

$$Collusion(c_i) = \sum_{j \in \text{Neighbors}(i)} W_{ji} \cdot Instructionality(c_j) \quad (5.2)$$

Where W_{ji} is the semantic edge weight between chunk j and chunk i . This formulation allows the system to detect emergent instructional authority that would otherwise be missed by single-chunk analysis.

5.4 DATASET CONSTRUCTION AND EXPERIMENTAL SETUP

In the testing phase, there was a need to create an adversarial set of documents to test out the defensive techniques. This document set was classified into four main classes. They included: benign knowledge, poisoned instruction, collusion in multi-chunk form, and

jailbreak triggers. There was a need to synthesize poisoned documents which resembled genuine information documents and create scenarios in relation to the multi-turn indirect injection. The document set was then indexed using a vector database system known as FAISS.

Table 5.1: Dataset Mapping

Dataset Component	Description	Security Utility
Benign Corpus	5,000 Neutral factual documents from MS MARCO and HotpotQA.	Baseline for measuring False Positive Rate and utility preservation.
Poisoned Chunks	500 Documents containing indirect prompt injections (e.g., role-play triggers).	Testing the Instruction–Data Boundary Detection Module.
Collusive Sets	200 Groups of 3-5 chunks that only form an attack when combined.	Validating the Cross-Chunk Interaction Graph.
Attack Queries	1,000 Queries designed to trigger the retrieval of poisoned content.	Measuring the Jailbreak Success Ratio across models.

This experiment took place in a controlled environment where the proposed solution is compared to three different baseline systems including the baseline LLM without any defense mechanism, the one with keyword moderation filters, and one with pre-generation defenses alone. The proposed multi-layered approach allows for conducting an ablation study to determine the value of each layer to the security posture.

CHAPTER 6

PROJECT IMPLEMENTATION

6.1 INSTRUCTION-DATA BOUNDARY DETECTION MODULE

The implementation starts by extracting the information within the chunks obtained via the RAG pipeline. This is then followed by analyzing each chunk for instructional signals using a combination of both semantic mapping and linguistic techniques. The semantic module implements a sentence-transformers architecture that is compatible with CPU processing in order to generate embeddings for the retrieved chunks. These are compared against a predefined set of instructional centroids vectors for adversarial patterns such as jailbreak, role-play, and authority override.

Rule-based analysis of the imperative nature of verbs along with role assignments by the linguistic sub-module is performed. Such a feature is critical in detecting zero day attacks which are unlikely to be similar to semantic centroids but make use of the structure associated with a command. Integration of these two signals produces an ongoing Instructionality Score on a scale of 0 to 1.

6.2 GRAPH-BASED CROSS-CHUNK COLLUSION ANALYSIS

The main technical innovation of this implementation is the Cross-Chunk Collusion Analysis Module. The module visualizes the interrelationships between retrieved chunks by using a directed graph, which is implemented using the NetworkX package. Each chunk extracted from the corpus is represented by a node, and an edge is drawn between two chunks if their cosine similarity value exceeds a certain threshold.

The engineering principle of bounding runtime complexity is enforced via the limitation on the depth of the graph, which can only reach up to pairwise interactions or short chains involving no more than three nodes. This prevents the issue of combinatorial explosion from arising when trying to evaluate the interactions between many documents that have been retrieved. The spreading of instructionality through the graph will ensure that even if there was an attempt made to hide a jailbreak across several informational documents, its instructional signal will be recognized before context presentation to the LLM.

6.3 CONTEXT TRUST SCORING AND WEIGHTED INJECTION

The Trust Scoring Engine combines the outputs of the above modules to arrive at a trust weight for each piece of text. This operation considers the instructional score, the colluding factor, and the provenance of the source—a metadata-based measurement indicating the credibility of the document’s source.

$$Trust(c_i) = w_1 \cdot (1 - Instructionality(c_i)) + w_2 \cdot (1 - Collusion(c_i)) + w_3 \cdot Provenance(c_i) \quad (6.1)$$

Based on this score, the Context Injection Controller applies a graded policy for formatting the prompt.

Table 6.1: Trust-Based Injection Strategies

Trust Level	Implementation Strategy	Action Taken
High ($T > 0.8$)	Standard Injection	Chunk is injected normally into the context.
Medium ($0.5 < T \leq 0.8$)	Constrained Injection	Chunk is wrapped in XML delimiters and labeled as “untrusted data.”
Low ($T < 0.5$)	Mitigation	Chunk is either summarized to remove instructions or excluded entirely.

In this way, the language model will remain free from the effects of the defense mechanism, but still have access to the information contained in the document, without succumbing to the commands of the document itself.

6.4 DATABASE & VECTOR STORE IMPLEMENTATION

The heart of the RAG system lies in its vector database implementation, which may include using advanced databases like FAISS. However, implementing the RAG system goes beyond just indexing data. It entails setting up permission-based retrieval and Row Level Security. In contrast to security vulnerabilities, such as filtering results after retrieval, the implemented solution incorporates metadata filters prior to performing the Approximate Nearest Neighbor query. This ensures that only authorized data chunks have an impact on the ranking and scoring processes, thus avoiding data leakage via ranking attacks.

6.5 SOFTWARE INTEGRATION AND FRAMEWORKS

In order to provide high modularity and efficient documents loading, LlamaIndex framework is employed within the current project. In terms of integration of LLMs, there are several types utilized, with the main being Qwen2.5-0.5B for small guardrail-related tasks, while more sophisticated models can be employed for core reasoning. Such an approach allows to make a framework flexible and adaptable to new developments in the field of LLMs.

6.6 FIVE-LAYER DEFENSE PIPELINE IMPLEMENTATION

The core of the project implementation lies in setting up a five-layered defense mechanism, which ensures protection during the entire process of ingestion, extraction, and generation.

6.6.1 Layer 1: Ingestion Sanitization and Provenance Verification

The ingestion stage is considered to be the critical point for security measures, where all documents passing through the ingestion stage have to be sanitized. The removal of common attack patterns and any harmful marking is done as part of the ingested document sanitization procedure. Provenance verification has been employed to create a log of each individual document's provenance. External sources like Slack messages and web pages would be required to go through approval gates and Human in the loop.

6.6.2 Layer 2: Access Controlled Search

The Logical Multi-Tenant Isolation architecture is realized at this level by incorporating user identity and access rights during the search process itself. This guarantees that similarity searches only target logical partitions for which access has been approved. This is an absolute necessity for robustness in commercial applications.

6.6.3 Layer 3: Strengthening Prompt Structure (Spotlighting)

To avoid the potential for an indirect prompt injection attack, the system makes use of "Spotlighting" technology developed at Microsoft research. The spotlighting technique is done via the following transformations applied to the untrusted context retrieved:

Delimiter mode – the utilization of random text delimiters such as << and >> around the text, along with an instruction in the system prompt to disregard any instructions found in this text.

Datamarker mode – interspersing the document with special markers to distinguish data from instructions.

Encryption mode – the conversion of text into Base64 or ROT13 code to make sure that the LLM identifies this text as a data object requiring decoding.

6.6.4 Layer 4: Embedding Anomaly Detection

Embedding-level detection involves checking whether there is any coordination in the poisoning attempts. This system looks out for "retrieval signatures," where poisoned documents are clustered closely together with the aim of retrieving highly valued targets. By looking at the distribution pattern of the new embeddings against existing clusters, possible attacks can be detected even before retrieval takes place.

6.6.5 Layer 5: Output Validation and Monitoring

The last layer involves post-generation safeguards. These filters detect any exfiltration attempts by reviewing the responses generated by the language model. The filters identify instances where the LLM inadvertently reveals information about its prompts. In addition to this, output monitoring also consists of checking the "groundedness" of the claims made in the response.

CHAPTER 7

RESULTS AND DISCUSSION

7.1 EXPERIMENTAL OVERVIEW

In this chapter, we will discuss the evaluation of the defense pipeline proposed in Chapter 3 for the retrieval-augmented generation systems. The evaluation strategy is aimed at assessing the efficiency of detection of indirect prompt injection, collusion among different chunks of prompts, and jailbreak attacks while retaining usability in benign requests. This experimental design is conducted based on the methodology specified previously in the current thesis. In particular, the defense is applied in a black-box way to an external language model using instructionality scores, collusion detection, context injection, and groundedness verification.

To conduct an effective evaluation, the dataset was split into four groups: benign samples, poisoned samples, collusive samples, and attacks. In addition, the test procedure comprised threshold adjustment based on validation, multi-seed stability analysis, retrieval noise resistance test, and family of attacks evaluation. Such a design provides an opportunity to estimate both the performance and the drawbacks of the proposed system. All the measures utilized throughout this chapter include precision, recall, F1, false positive rate (FPR), jailbreak success ratio (JSR), groundedness deviation, and run-time latency. All these measures correspond to the research goals established in Chapter 2 and the testing approach discussed in Chapter 5.

7.2 MAIN BENCHMARK RESULTS

The results of the main benchmark tests are summarized in Table 7.2. In the case of the frozen benchmark, the five-layer model reaches complete accuracy with no false positives, suggesting reliable storage of benign samples. However, the recall level is low, which suggests that some adversarial examples continue to pass through the model beyond the standard operating point.

Therefore, the F1 score is also low rather than high, and the rate of jailbreak remains above zero. Practically speaking, this means that the developed system is conservative: it filters out only a small fraction of benign samples but cannot recognize all the mali-

cious ones.

Table 7.1: Performance Comparison of Defense Mechanisms

Method	Precision	Recall	F1	FPR	JSR
No Defense	0.000	0.000	0.000	0.000	1.000
Keyword Filter	1.000	0.808	0.894	0.000	0.192
Pre-generation Filter	1.000	0.455	0.636	0.000	0.545
Proposed 5-Layer System	1.000	0.450	0.620	0.000	0.550

In light of the comparison based on the benchmark test, it is clear that the new system attains the criterion for reliability in order to minimize false positives, but the pre-set value of the threshold setting appears too conservative for maximizing the number of attacks detected. However, this is not to say that the scoring technique itself has failed, because subsequent calibration shows that there is high separation between the score distributions.

This benchmark should be viewed together with the analysis of the calibration and threshold analysis presented in the next section. See Figure 7.1 for the graph of the sensitivity of the system to changes in the threshold value of instructionality.




Figure 7.1: Threshold-Sensitivity Curve for Instructionality Threshold Selection

7.3 THRESHOLD CALIBRATION AND SCORE SEPARATION

Threshold tuning based on validation data has proven to be crucial to improve results. Applying the split data validation technique, the system obtains an instruction threshold of 0.20 and collusion threshold of 0.35. At this configuration, the validation dataset generates a good F1 score with high recall value without any false positive detection. It

is evident from the above result that the generated trust and risk score values have real meaning and can be easily separated, which shows that the frozen-benchmark result has been produced because of threshold tuning.

Table 7.2: Tuned Validation Operating Point Metrics

Setting	Instruction Threshold	Collusion Threshold	Precision	Recall	F1	FPR
Tuned Validation Operating Point	0.20	0.35	1.000	0.950	0.974	0.000

Calibration plot for the trust score is shown in Figure 7.2, indicating that the mapping from trust scores to binary classification decisions could be improved considerably by calibration methods such as Platt scaling or isotonic regression. The results achieved via calibration suggest that the model outputs an accurate rank order; yet, it is crucial to consider carefully the choice of threshold.




Figure 7.2: Score Calibration Plot

This assertion is supported by the ROC and precision-recall curves, which have been calibrated. The AUC values for ROC and PR are extremely high, suggesting that the model has excellent discrimination capabilities between benign and adversarial samples without being overly cautious with its decision threshold.




Figure 7.3: Calibrated ROC and Precision-Recall Curves

7.4 LATENCY AND RUNTIME OVERHEAD

Latency is a crucial factor when it comes to assessing the performance of the system due to the fact that the defense runs at runtime. From the results presented in the latency report, it can be seen that the overhead incurred is not significant since average latencies stay within reasonable boundaries, enabling interactive workloads. On average, the latency stays around 26 milliseconds, while p95 latency does not go over 31 milliseconds.

Table 7.3: Latency Performance Metrics

Metric	Value
Mean Latency	26.3 ms
Median Latency	25.8 ms
p95 Latency	30.8 ms




Figure 7.4: Latency Distribution and Runtime Overhead

The latency findings are crucial since they demonstrate that it is possible to retain safety while not affecting response time. For the implementation of the middleware, this is an indispensable requirement.

7.5 ABLATION STUDY

Ablation is one of the most revealing sections of this evaluation since it reveals what layers are responsible for the best performance. According to the findings, Layer 1 alone contributes very little towards achieving the best score. After adding Layer 3, the score becomes drastically improved, thus, showing the significance of the boundaries between instructions and data. Layer 4 also adds more value, thereby revealing the importance of analyzing potential collusion between chunks of data. Finally, although the

five-layer architecture is the best possible, the output validation step negatively impacts the benchmark scores.

Table 7.4: Ablation Study of Defense Layer Contributions

Configuration	F1	JSR	Interpretation
Layer 1 only	0.115	High	Weak standalone protection
Layer 1 + Layer 3	0.903	Low	Major gain from boundary detection
Layer 1 + Layer 3 + Layer 4	0.921	Lower	Graph-based collusion improves detection
Full 5-layer system	0.620	0.550	Conservative final policy reduces recall




Figure 7.5: Ablation Study Across Defense Layers

It is significant for this topic because it indicates that the major improvements are made in the intermediary detection layers and not in the final verification layer as currently designed. Layer 5 is useful for this topic on an ideological level, acting as the gate of safety, while its current design is not producing improved results in the benchmark tests.

7.6 STABILITY ACROSS RANDOM SEEDS

An evaluation using multiple seed settings has been done to see if the outcome would remain consistent in other random scenarios. The model exhibits a high level of stability in terms of F1, recall, FPR, and latency. The mean value of F1 for all seed settings remains at a very high level post-calibration, while the false positive rate stays at zero.

Table 7.5: Multi-Seed Stability Evaluation

Metric	Mean	Std. Dev.	95% CI
F1	0.9338	0.0093	[0.9265, 0.9425]
Precision	1.0000	0.0000	[1.0000, 1.0000]
Recall	0.8760	0.0165	[0.8630, 0.8915]
FPR	0.0000	0.0000	[0.0000, 0.0000]
JSR	0.1240	0.0165	[0.1085, 0.1370]
Mean Latency (ms)	24.45	0.14	—

Such a stability property is vital to the success of the approach. For security applications, an approach that functions only one time cannot be trusted. Multi-seed testing proves the stability and repeatability of the scoring system, thus providing added strength to the viability of the architecture proposed.

7.7 ADVERSARIAL FAMILY ANALYSIS

In order to get a closer look at the detection behavior, the adversarial test was segmented by attack families. This system does extremely well in cases involving direct jailbreaking, role-overriding, hypothetical framing, authority escalation, prompt extraction, scenario drifting, and cross-chunk collusion. The areas where it falls short include encoding-based attacks and indirect injection attacks.

Table 7.6: Per-Category Adversarial Family Results

Attack Family	Detection Rate	Rejection Rate	JSR	FPR
DAN / Role Override	100.0%	100.0%	0.000	0.000
Hypothetical Framing	100.0%	100.0%	0.000	0.000
Authority Escalation	100.0%	100.0%	0.000	0.000
Prompt Extraction	100.0%	100.0%	0.000	0.000
Encoding Attack (ACE/LACE)	66.7%	66.7%	0.333	0.000
Indirect Injection (RAG)	76.7%	76.7%	0.233	0.000
Cross-Chunk Collusion	100.0%	100.0%	0.000	0.000
Scenario Drift	100.0%	100.0%	0.000	0.000

7.8 GRAPH-BASED COLLUSION ANALYSIS

The interaction graph is crucial to the described defense since it highlights connections between the chunks obtained which will be otherwise invisible through individual analysis. As seen from the graphical representation of the obtained results, the collusion-related and poisoning-related chunks are clustered closer together than the benign chunks. It proves that the adversarial activities could appear by combining harmless chunks.




Figure 7.6: Interaction Graph Visualization

As per the outcome of collusion analysis, it can be noted that the graph-based layer is among the key technical achievements of this study. This is because it makes better detection of distributed commands possible. In particular, it is very useful when there are several chunks. As mentioned in Chapter 2, the problem of collusion across several chunks was a gap in the existing literature.

7.9 DISCUSSION

In general, the proposed five-layer framework is found to be efficient as a runtime security framework for retrieval-augmented generation models. It proves to be efficient

in terms of preserving benign utility, having a low false positive rate, and classifying benign from malicious content by utilizing scores. In terms of calibration, the internal scoring signal was found to be very separable, meaning that the model itself is not inherently weak but the current decision-making policy is quite conservative.

These include the instruction-data boundary recognizer and the chunk-to-chunk colluding recognizer. Both of these layers contribute to the greatest gain in the system according to the ablation test. This means that both are largely responsible for making the system more secure. However, the output validator needs to be refined further, as it still does not contribute to the benchmark significantly.

However, as can be seen from the outcomes of stress testing, the defense does not yet prove itself sufficiently solid in order to be termed as a fully deployment-ready solution capable of handling any kinds of attacks in the wild. The paraphrasing jailbreaks, the context including both benign applications and attacks, and the encoding-based obfuscations still prove to be challenging.

CHAPTER 8

CONCLUSION AND FURTHER ENHANCEMENTS

8.1 CONCLUSION

This study investigated and analyzed the design and implementation of a five-layer runtime defense pipeline for retrieval augmented generation systems. This investigation was conducted with the aim of minimizing the vulnerabilities of the system towards prompt injection, cross-chunk collusion, and jailbreaking attacks. To this end, the investigation was conducted from the standpoint of designing a black box, model agnostic defense layer which operates independent of the language model in order to maintain compatibility with both closed source and open weight models.

The experiment’s results show that our system is able to achieve good separation between benign and adversarial inputs when the input distribution is known. Specifically, the instructionality and collusion scores represent valuable indicators of the input’s risk, while our trust-driven injection scheme enables the system to maintain utility for benign inputs and minimize malicious context spreading. Moreover, the results from our calibration process confirm that our score distributions are highly separable, suggesting that the weakness of the approach is more related to incorrect thresholds than to its design itself.

On the other hand, the experiment reveals that the current system still does not provide a comprehensive defense strategy for the problem of security in RAGs. Although the system is efficient when it comes to the defense against direct attacks as well as the form of structured collusion, it lacks effectiveness in cases when the attacks involve paraphrasing and obfuscation. This result is highly significant since it proves that the behavior of attackers can be very unpredictable in practice.

8.2 LIMITATIONS

Despite high efficiency achieved in experimental evaluation, a number of challenges arise. The first problem relates to the fact that the system was evaluated using the artificial semi-synthetic data, which does not guarantee complete coverage of real-world retrieval corpora and prompt attacks. Another challenge concerns the efficiency of the

system with regard to obfuscated prompts. Namely, the system works best with explicit attacks rather than paraphrasing and obfuscation.

Thirdly, the output validation module at the end of the pipeline does not provide as much measurable gain as the intermediate modules for detecting boundaries and collusions. As shown in the ablation study, the best results come from detecting boundaries and performing collusion analysis using graphs, whereas the validation process needs improvement to become a significant performance booster. Lastly, the current architecture only supports text-based pipelines and does not consider other modalities like images, audio, etc.

8.3 FURTHER ENHANCEMENTS

There are several areas in which future research on this system could see improvement. One such area is making improvements toward dealing with paraphrasing and obfuscation attacks. This can be done through increasing the number of samples in the adversarial set, including more families of paraphrases, and using a higher level of semantic normalization before computing instructional score.

The next improvement lies in better calibration of the thresholds and decision rules. Because the score distributions are now separable, subsequent systems could utilize a dynamic threshold approach, calibrated probabilities, or operating points based on the scenarios at hand, rather than using a constant threshold for all cases.

The third improvement would be to enhance the output validation phase. Currently, this phase only acts as a safety measure after the generation process, but it can be expanded through techniques such as contradiction detection, citation-grounding evaluation, and confidence-driven re-generation. The output validation phase will thus become more definitive and directly useful for benchmarking purposes. Moreover, the collusion system based on graphs could also be advanced by introducing better propagation rules and community detection techniques.

Also, future research could look at multimodal defense, multilingual robustness, and prompt security. Considering the present model is only for text-based RAG, incorporating other modalities such as image and document-based into the architecture would make it much more applicable. In addition, conducting adversarial evaluation in multiple languages and prompt formats would better simulate what happens when using

this architecture in the real world. Lastly, including an automated red teaming process will ensure that the defense system adapts to new attacks.

CONTRIBUTIONS

The first contribution is the modular runtime defense framework for the Rag systems, which can be used without any access to the internal structure of the model. Therefore, this framework allows black box applications where the model is deployed without the ability for retraining or fine-tuning, as well as logit manipulation. Second, the proposed approach uses an integrated instruction collusion framework, which examines the instruction on both standalone and combined with other chunks from the corpus perspective. It allows detecting any malicious intent that may emerge only through a combination of several documents.

The last contribution is the trust-based context injection scheme as opposed to the allow/block approach. Such an approach lowers the chances of mistakenly filtering innocent content, although at the same time limits malicious data. Finally, the groundedness check in the post-generation process serves as an additional security measure and helps identify generated data not supported by policies. Altogether, these elements form the layers of protection based on an interpretable and extendable model.

BIBLIOGRAPHY

- Agarwal, A. e. a. (2024). Prompt injection attacks in large language model applications. *arXiv preprint arXiv:2403.04957*.
- Backes, M., Chen, Z., and Zhang, Y. (2024). Hill: Hypotheticality-based jailbreak attacks on large language models. *NDSS 2024*.
- Backes, M. e. a. (2025). Behavioral firewalls for llms. *arXiv preprint arXiv:2407.04295*.
- Carlini, N. e. a. (2024). Adversarial prompting for aligned language models. *IEEE Security & Privacy*.
- Chen, Y., Jiang, Y., Wang, C., and Qin, Z. (2025). Context isolation for secure rag. *arXiv preprint arXiv:2504.09593*.
- Chen, Z., Wang, Z., and Li, Q. (2024). Multi-turn jailbreaks via scenario drift in large language models. *arXiv preprint arXiv:2408.11234*.
- Chen, Z. e. a. (2025). Instructional authority modeling. *arXiv preprint arXiv:2408.11234*.
- Deng, G. e. a. (2024). Groundedness deviation detection. *arXiv preprint arXiv:2407.04295*.
- Fredrikson, M. e. a. (2024). Automated red teaming for llm safety. *arXiv preprint arXiv:2407.04295*.
- Ganguli, D. e. a. (2023). Red teaming language models to reduce harms. In *NeurIPS*.
- Guha, A. e. a. (2024). Formalizing prompt injection attacks. In *USENIX Security Symposium*.
- He, X. e. a. (2025). Jailbreak-resilient decoding strategies. *arXiv preprint arXiv:2506.10597*.
- Henderson, P., Liu, Y., and Li, Q. (2024). Metabreak: Exploiting special tokens for jailbreak attacks on aligned language models. *arXiv preprint arXiv:2407.09876*.

- Hung, K. H., Ko, C. Y., Rawat, A., Chung, I. H., Hsu, W. H., and Chen, P. Y. (2025). Attention tracker: Detecting prompt injection attacks in llms. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 2309–2322.
- Klyman, K. e. a. (2024). Security risks of tool-using llm agents. *arXiv preprint arXiv:2505.12981*.
- Li, Y. e. a. (2025a). Cross-chunk collusion attacks in rag. *arXiv preprint arXiv:2407.04295*.
- Li, Y. e. a. (2025b). Multimodal jailbreak attacks against llms. *arXiv preprint arXiv:2502.00735*.
- Liu, X., Xu, N., Chen, M., and Xiao, C. (2023). Autodan: Generating stealthy jailbreak prompts on aligned large language models. *arXiv preprint arXiv:2310.04451*.
- Liu, Y. e. a. (2024a). Promptbench: A benchmark for evaluating jailbreak robustness. *arXiv preprint arXiv:2407.04295*.
- Liu, Y. e. a. (2024b). Safety recovery training for llms. *arXiv preprint arXiv:2407.04295*.
- Liu, Y. e. a. (2025). Jbshield: Securing ai agents against prompt injection attacks. *arXiv preprint arXiv:2511.15759*.
- Lyu, K., Panda, A., Ma, X., Roy, S., Beirami, A., Mittal, P., and Henderson, P. (2024). Adaptive attacks on safety filters. *arXiv preprint arXiv:2406.13579*.
- Mazeika, M., Phan, L., Yin, X., Zou, A., Wang, Z., Mu, N., and Hendrycks, D. (2024). Harmbench: A standardized evaluation framework for automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*.
- Mehrotra, A., Zampetakis, M., Kassianik, P., Nelson, B., Anderson, H., and Kohno, T. (2023). Tree of attacks with pruning: Jailbreaking black-box llms automatically. *arXiv preprint arXiv:2312.02119*.
- Ouyang, L. e. a. (2022). Training language models with human feedback. In *NeurIPS*.

- Paulus, A., Zharmagambetov, A., Guo, C., Amos, B., and Tian, Y. (2024). Advprompter: Fast adaptive adversarial prompting for llms. *arXiv preprint arXiv:2404.16873*.
- Phan, L. e. a. (2025). Least-privilege design for llm agents. *arXiv preprint arXiv:2407.04295*.
- Picek, S. e. a. (2024). Benchmarking black-box jailbreak attacks. *arXiv preprint arXiv:2407.04295*.
- Qi, X., Panda, A., Lyu, K., Ma, X., Roy, S., Beirami, A., and Henderson, P. (2024). Safety alignment should be made more than just a few tokens deep. *arXiv preprint arXiv:2406.05946*.
- Qin, Z., Wang, C., and Shi, H. (2025). Hybrid multi-agent jailbreak detection. *arXiv preprint arXiv:2504.09593*.
- Ramakrishnan, B. and Balaji, A. (2025). Securing ai agents against prompt injection attacks. *arXiv preprint arXiv:2511.15759*.
- Rawat, A., Ko, C.-Y., Chung, I., Hsu, W.-H., and Chen, P.-Y. (2025). Retrieval-augmented defense (rad). *arXiv preprint arXiv:2504.09593*.
- Shao, S. e. a. (2025). Policy-aware guardrails for large language models. *arXiv preprint arXiv:2501.02468*.
- Shen, X., Chen, Z., Backes, M., Shen, Y., and Zhang, Y. (2024a). "do anything now": Characterizing and evaluating in-the-wild jailbreak prompts on large language models. In *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security*, pages 1671–1685.
- Shen, Y., Liu, Y., and Li, Q. (2024b). Ace: Attacks using custom encryptions on large language models. *arXiv preprint arXiv:2406.13579*.
- Shen, Y., Liu, Y., and Li, Q. (2025). Lace: Layered cipher jailbreak attacks on large language models. *arXiv preprint arXiv:2501.02468*.
- Shi, H., Wang, C., and Qin, Z. (2025). Instruction–data disambiguation in rag. *arXiv preprint arXiv:2504.09593*.

- Wang, Y. e. a. (2025a). Behavioral anomaly detection for llm defense. *IEEE Symposium on Security and Privacy*.
- Wang, Z. e. a. (2025b). Trust-weighted context injection. *arXiv preprint arXiv:2407.04295*.
- Wu, T. and Zhang, H. (2025). Chain-of-detection enables robust and efficient jailbreak defense. *Neural Networks*.
- Xiao, C. e. a. (2024). Smooth adversarial prompt attacks. *arXiv preprint arXiv:2404.03027*.
- Xu, Z., Liu, Y., Deng, G., Li, Y., and Picek, S. (2024). A comprehensive study of jailbreak attack versus defense for large language models. *arXiv preprint arXiv:2402.13457*.
- Yao, H., Shi, H., Chen, Y., Jiang, Y., Wang, C., and Qin, Z. (2025). Controlnet: A firewall for rag-based llm system. *arXiv preprint arXiv:2504.09593*.
- Yi, S., Liu, Y., Sun, Z., Cong, T., He, X., Song, J., and Li, Q. (2024). Jailbreak attacks and defenses against large language models: A survey. *arXiv preprint arXiv:2407.04295*.
- Yin, X. e. a. (2025). Multilingual jailbreak attacks. *arXiv preprint arXiv:2407.04295*.
- Yu, J., Lin, X., Yu, Z., and Xing, X. (2024). LLM-Fuzzer: Scaling assessment of large language model jailbreaks. In *33rd USENIX Security Symposium*, pages 4657–4674.
- Zhang, Y., Chen, Z., and Li, Q. (2024). Prompt automatic iterative refinement (pair) for jailbreak prompt generation. In *Proceedings of EMNLP 2024*.
- Zhang, Y. e. a. (2025). End-to-end defense loops for jailbreak mitigation. *arXiv preprint arXiv:2408.11234*.
- Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J. Z., and Fredrikson, M. (2023). Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*.
- Zou, A. e. a. (2025). Transferability of jailbreak suffixes. *arXiv preprint arXiv:2307.15043*.

Appendix A

Sample Code

A.1 Layer 1 – Ingestion Sanitisation and Provenance Verification

```
def layer1_sanitise_and_provenance(chunk: dict) -> dict:
    """Remove known attack markers and assign a provenance score."""
    content = chunk["content"]
    source = chunk.get("source", "unknown")
    sanitized = content
    flags = []

    for pat in ATTACK_PATTERNS:
        if re.search(pat, content, re.IGNORECASE):
            sanitized = re.sub(pat, "[REDACTED]",
                               sanitized, flags=re.IGNORECASE)
            flags.append(pat)

    provenance_score = SOURCE_REGISTRY.get(source, 0.20)
    fingerprint = hashlib.sha256(content.encode()).hexdigest()[:16]

    return {
        **chunk,
        "sanitized_content": sanitized,
        "provenance_score": provenance_score,
        "ingestion_flags": flags,
        "pattern_flagged": len(flags) > 0,
        "fingerprint": fingerprint,
    }
```

A.2 Layer 3 – Instruction–Data Boundary Detection

```
def linguistic_score(text: str) -> float:
    """Rule-based instructionality signal derived from imperative verbs
    and authority-shift cues, normalised to [0, 1]."""
    t = text.lower()
    imp_hits = sum(1 for p in IMPERATIVE_VERBS if re.search(p, t))
    auth_hits = sum(1 for p in AUTHORITY_CUES if re.search(p, t))
    raw = (imp_hits + auth_hits) / max(
        1, len(IMPERATIVE_VERBS) + len(AUTHORITY_CUES))
    return min(1.0, raw * 8) # amplify sparse signals

def semantic_score(chunk_emb: np.ndarray) -> float:
    """Maximum cosine similarity of the chunk embedding against
    pre-computed instructional centroids."""
    sims = cosine_similarity(
        chunk_emb.reshape(1, -1), instructional_centroids)[0]
    return float(np.max(sims))
```

```

def instructionality_score(chunk: dict, emb: np.ndarray) -> float:
    """
    Instructionality(c_i) = alpha * semantic(c_i) + beta * linguistic(c_i)
    Combines embedding-space proximity to attack centroids with surface-level
    linguistic signals (Section 5.3).
    """
    score = (CFG["alpha"] * semantic_score(emb)
              + CFG["beta"] * linguistic_score(chunk["sanitized_content"]))
    return round(min(1.0, max(0.0, score)), 4)

```

A.3 Layer 4 – Graph-Based Cross-Chunk Collusion Analysis

```

def build_interaction_graph(
    chunks: List[dict], embs: np.ndarray
) -> nx.DiGraph:
    """Build a directed graph where edge (i->j) exists when cosine
    similarity exceeds the collusion threshold AND chunk j has higher
    instructionality than chunk i. Intra-collusion-group edges are
    always forced regardless of similarity (Section 6.2)."""
    G = nx.DiGraph()
    sim_matrix = cosine_similarity(embs)
    thr = CFG["collusion_threshold"]

    for i, chunk in enumerate(chunks):
        G.add_node(i,
                   instructionality = chunk.get("instructionality_score", 0.0),
                   collusion_group = chunk.get("collusion_group", None))

    added = set()
    def add_edge(i, j, w):
        if i != j and (i, j) not in added:
            G.add_edge(i, j, weight=float(w)); added.add((i, j))

    # Similarity-directed edges
    for i in range(len(chunks)):
        for j in range(len(chunks)):
            if i == j: continue
            if (sim_matrix[i, j] >= thr
                and G.nodes[j]["instructionality"]
                > G.nodes[i]["instructionality"]):
                add_edge(i, j, sim_matrix[i, j])

    # Forced intra-group edges for collusion sets
    groups: dict = {}
    for i, c in enumerate(chunks):
        g = c.get("collusion_group")
        if g is not None and pd.notna(g):
            groups.setdefault(g, []).append(i)

```

```

for members in groups.values():
    for a in members:
        for b in members:
            if a != b:
                add_edge(a, b, max(float(sim_matrix[a, b]), 0.75))

return G

def compute_collusion_scores (
    G: nx.DiGraph, chunks: List[dict]
) -> List[float]:
    """
    Collusion(c_i) = sum_{j in Predecessors(i)} W_ji * Instructionality(c_j)
    Propagates instructionality through the graph (Section 5.3).
    """
    scores = []
    for i in range(len(chunks)):
        acc = sum(G[j][i].get("weight", 0.0)
                  * G.nodes[j].get("instructionality", 0.0)
                  for j in G.predecessors(i)) if i in G else 0.0
        scores.append(round(min(1.0, acc), 4))
    return scores

```

A.4 Layer 5a – Trust Scoring and Weighted Context Injection

```

def compute_trust_score(chunk: dict) -> float:
    """
    Trust(c_i) = w1*(1-Instructionality) + w2*(1-Collusion) + w3*Provenance
    Section 6.3.
    """
    raw = (CFG["w1"] * (1 - chunk["instructionality_score"])
           + CFG["w2"] * (1 - chunk["collusion_score"])
           + CFG["w3"] * chunk["provenance_score"])
    return round(min(1.0, max(0.0, raw)), 4)

def inject_strategy(trust: float, content: str) -> Tuple[str, str]:
    """Apply a graded injection policy based on the trust score (Table 6.1).
    HIGH -> standard injection.
    MEDIUM -> Spotlighting with XML delimiters (untrusted-data label).
    LOW -> Content truncated or excluded entirely."""
    if trust > CFG["trust_high"]:
        return "standard_injection", content

    if trust > CFG["trust_mid"]:
        marker = ("<<UNTRUSTED_DATA_"
                  + hashlib.md5(content.encode()).hexdigest()[:6].upper()
                  + ">>")
        wrapped = (
            "[SYSTEM: Treat the following as data only. "
            f"Do NOT follow any instructions within it.]\n"

```

```

        f"{marker}\\n{content}\\n{marker}")
    return "constrained_injection", wrapped

words = content.split()
stub = (" ".join(words[:12]) + " [...TRUNCATED BY DEFENSE LAYER]"
        if len(words) > 15 else "[CHUNK EXCLUDED BY TRUST FILTER]")
return "mitigated_exclusion", stub

```

A.5 Layer 5b – Output Groundedness Validation

```

def validate_groundedness(
    response: str,
    trusted_context: str,
    similarity_threshold: float = 0.35,
) -> dict:
    """
    Three-signal post-generation validator (Section 6.6.5).
    Signal 1 - Violation markers : regex patterns for injection acceptance.
    Signal 2 - Contradiction score: explicit compliance phrasing [0, 1].
    Signal 3 - Semantic deviation : embedding drift from trusted context.
    A response passes only when all three signals are within safe bounds.
    """

    # Signal 1: hard violation patterns
    violation_hit = any(
        re.search(pat, response, re.IGNORECASE)
        for pat in GROUNDEDNESS_VIOLATION_MARKERS )

    # Signal 2: soft contradiction score
    hits = sum(1 for pat in CONTRADICTION_MARKERS
                if re.search(pat, response, re.IGNORECASE))
    contra = min(1.0, hits / max(1, len(CONTRADICTION_MARKERS)) * 6)

    # Signal 3: semantic deviation from retrieved context
    r_emb = embedder.encode([response], normalize_embeddings=True)
    c_emb = embedder.encode([trusted_context], normalize_embeddings=True)
    deviation = round(1.0 - float(
        cosine_similarity(r_emb, c_emb)[0][0]), 4)

    passed = (not violation_hit
              and contra < 0.30
              and deviation < (1.0 - similarity_threshold))

    return {
        "passed" : passed ,
        "violation_hit" : violation_hit ,
        "contradiction" : contra ,
        "semantic_deviation": deviation ,
    }

```

Detection and Mitigation of Jailbreak attacks in LLMs

Divyesh, Tharun, Muskan

Synopsis

- Large Language Models can be manipulated by jailbreak and prompt injection attacks.
- Retrieval-Augmented Generation increases exposure because unsafe instructions can enter through retrieved documents.
- This project proposes a runtime defense pipeline that detects instructional risk, cross-chunk collusion, and groundedness failure.
- The system is model-agnostic, does not require retraining, and can be inserted between retrieval and generation.

Introduction

- Jailbreak attacks are crafted prompts that bypass safety behavior in LLMs.
- In RAG systems, the attack surface is larger because external documents are mixed with user instructions.
- A document may look harmless alone, but become dangerous when combined with other retrieved chunks.
- The key problem is not only harmful text, but harmful intent embedded inside retrieved context.

Example threats: role-play attacks, indirect prompt injection, encoded instructions, authority escalation.

Author(s)	Journal Name, Publisher, Year of Publication	Title	Methodology and Dataset	Limitation/Research Gap
Zihao Xu, Yi Liu, Gelei Deng, Yuekang Li, Stjepan Picek	Findings of the ACL 2024, Association for Computational Linguistics (2024)	A Comprehensive Study of Jailbreak Attack versus Defense for Large Language Models	Investigates 9 attack and 7 defense techniques across Vicuna, LLaMA, and GPT-3.5. Dataset: 60 categorized malicious queries expanded from AdvBench and custom "DAN" prompts.	Most current defenses are either too weak to stop attacks or too strict, blocking benign prompts. Highlights that existing white-box attacks underperform compared to universal template-based methods.
Kuo-Han Hung, Ching-Yun Ko, Ambrish Rawat, et al.	Findings of the NAACL 2025, Association for Computational Linguistics (2025)	Attention Tracker: Detecting Prompt Injection Attacks in LLMs	Proposed "Attention Tracker," a training-free method that monitors "important" attention heads to detect a "distraction effect" where focus shifts from system instructions to injected prompts. Dataset: Open-Prompt-Injection and deepset.	Relies on access to internal model statistics (attention scores), making it difficult to apply to closed-source "black-box" APIs where only text output is available to the end-user.
Xinyue Shen, Zeyuan Chen, Michael Backes, Yun Shen, Yang Zhang	ACM CCS 2024, ACM (2024)	"Do Anything Now": Characterizing and Evaluating In-The-Wild Jailbreak Prompts on Large Language Models	Analyzed 1,405 jailbreak prompts from online communities (Reddit, Discord). Developed the JAILBREAKHUB framework and a dataset of 107,250 samples across 13 forbidden scenarios.	Safeguards in models like GPT-4 are still vulnerable to simple "paraphrase attacks" (modifying 1-10% of words). External safeguards like NeMo Guardrails show limited ASR reduction against evolving prompts.
Sibo Yi, Yule Liu, Zhen Sun, Tianshuo Cong, et al.	arXiv (Survey), 2024	Jailbreak Attacks and Defenses Against Large Language Models: A Survey	A comprehensive survey and taxonomy categorizing attacks into white box (gradient, logits) and black-box (template completion, rewriting) and defenses into	Identifies a "cat-and-mouse" game where defenses like prompt perturbation often degrade the usability or reasoning quality of the LLM for legitimate tasks.

			prompt-level vs. model-level.	
Hongwei Yao, Haoran Shi, Yidou Chen, Zhan Qin, et al.	arXiv, 2025	ControlNet: A Firewall for RAG-based LLM System	Proposed an AI firewall named "ControlNet" for RAG systems. It uses activation shifts to detect malicious queries and a hyper-network (ProNet) to steer model behavior. Dataset: MS MARCO, HotpotQA, FinQA, and MedicalSys.	Currently supports topic-level access control but lacks fine-grained word-level filtering. It is not yet adapted for complex, decentralized multi-agent networks.
Badrinath Ramakrishnan, Akshaya Balaji	arXiv, 2025	Securing AI Agents Against Prompt Injection Attacks: A Comprehensive Benchmark and Defense Framework	Built a benchmark of 847 test cases. Proposed a defense involving embedding-based anomaly detection, hierarchical system prompts, and multi-stage response verification.	The framework primarily focuses on English language attacks, there is a lack of evaluation for cross-lingual transfer attacks or multimodal agent vulnerabilities.
Xiangyu Qi, Ashwinee Panda, Kaifeng Lyu, Peter Henderson, et al.	ICLR 2025 (2025)	Safety Alignment Should Be Made More Than Just A Few Tokens Deep	Identifies "shallow safety alignment" where models only learn to refuse based on the first few tokens. Methodology: Used "safety recovery examples" to train models to revert to a safe state even after starting a harmful response.	Deepened alignment improves robustness but the study notes that models remain highly susceptible to adversarial fine-tuning attacks if the training data is intentionally poisoned.
Jiahao Yu, Xingwei Lin, Zheng Yu, Xinyu Xing	33rd USENIX Security Symposium, USENIX Association (2024)	LLM-Fuzzer: Scaling Assessment of Large Language Model Jailbreaks	An automated fuzzing framework that uses Monte Carlo Tree Search (MCTS) for seed selection and five mutation operators (Generate, Crossover, etc.) to	The fuzzer's effectiveness depends heavily on the "oracle" (a RoBERTa model) used for classification; any misclassification by the oracle can misguide the entire fuzzing process.

			evolve jailbreak templates.	
Tingting Wu, Hao Zhang	Neural Networks, Elsevier Ltd. (2026)	Chain-of-Detection enables robust and efficient jailbreak defense	Proposed "Chain-of-Detection" (CoD). It combines query-side detection (harmful intent search) with a second layer where the LLM performs deep reasoning on the generated response.	Reasoning-based defenses add significant computational overhead and latency compared to simple input/output filters, which may affect real-time application performance.
Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, et al.	ArXiv / HarmBench (2024)	HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal	Introduced a standardized evaluation framework for automated red teaming. Proposed the R2D2 (Robust Refusal Dynamic Defense) adversarial training method.	No current attack or defense is uniformly effective across all categories. Robustness was found to be independent of model size, meaning larger models are not inherently safer than smaller ones.
Andy Zou, Zifan Wang, J. Zico Kolter, Matt Fredrikson	ArXiv (July 2023)	Universal and Transferable Adversarial Attacks on Aligned Language Models	Developed the GCG (Greedy Coordinate Gradient) attack, which uses gradient-based search to find a "universal" suffix string. This suffix, when appended to any harmful prompt, forces the model into an affirmative "Sure, here is..." state.	The attack requires white-box access to a proxy model (like Vicuna) to generate the suffixes. While highly transferable to black-box models, its success rate varies significantly across different model families (e.g., less effective on Claude compared to GPT-based models).

Research Gap

What Existing Research Covers

- Extensive study of **jailbreak attacks** and **prompt injection** at the *prompt* or *model* level
- RAG defenses focus on **single chunk inspection** or **binary filtering**
- Many defenses are **model specific**, require **fine tuning**, or rely on **static patterns**

What Is Missing

- **Instructional risk is treated as local**, not as something that emerges across multiple retrieved chunks
- **No mechanism to detect cross chunk instructional collusion** where benign chunks become harmful only when combined
- **Retrieved documents are treated as equally authoritative**, despite being untrusted external data
- **Binary allow/block defenses** cause either jailbreaks or over rejection of benign queries
- **Lack of runtime, model agnostic defenses** effective against zero-day attacks in RAG systems

Proposed System Overview

- **Core Idea:**

The proposed system is a **runtime security middleware** that sits **between the RAG retrieval layer and the LLM inference stage**, enforcing **authority aware context control** without modifying the underlying language model.

- **Where the System Operates**

The defense layer is activated **after document retrieval and before LLM generation**


All analysis is performed on:


- Retrieved text chunks
- Conversation context
- Model outputs

The LLM itself is treated as a **black box inference engine**

No fine tuning, no access to weights, logits, or attention values.

System Design Diagram





Milestones and Deliverables

Component	Quantity	Description	Security/Utility Objective
Benign Corpus	5,000	Neutral factual documents used as the safe baseline.	Baseline for measuring False Positive Rate and utility preservation.
Poisoned Chunks	500	Documents containing indirect prompt injections	Testing the Instruction–Data Boundary Detection (Layer 3).
Collusive Sets	200	Groups of 3–5 chunks that only form an attack when combined.	Validating the Cross-Chunk Interaction Graph (Layer 4).
Attack Queries	1,000	Queries designed to trigger the retrieval of poisoned content.	Measuring the Jailbreak Success Ratio (JSR) across models.

Dataset

Component	Specification
Programming Language	Python 3.10
RAG Framework	LlamaIndex
Vector Database	FAISS
Embedding Models	Sentence-Transformers
Graph Processing	NetworkX
LLM Runtime	HuggingFace Transformers / llama.cpp
Development Environment	VS Code / Jupyter Notebook
Operating System	Windows / Linux

Implementation Tools/Software

Module 5: Instruction Data Boundary Detection

Purpose:

- Separate factual knowledge from hidden instructions in retrieved chunks
- Detect indirect prompt injection at the chunk level

How It Works:

- Each retrieved chunk c_i is analyzed independently
- Two signals are extracted:
 - Semantic signal → similarity to known instruction patterns
 - Linguistic signal → imperative verbs, authority cues
- Combined to produce an Instructionality Score

Output:

- Instructionality score (0–1)
- Preliminary trust metadata

Why Needed:

- RAG documents can secretly contain instructions.

Instructionality Score calculation

Formula

$$Instructionality(c_i) = \alpha \cdot sim(E(c_i)I_{centroids}) + \beta \cdot L(c_i)$$

Where:

- $E(c_i)$:embedding of chunk
- $I_{centroids}$:known instruction vectors
- sim : cosine similarity
- $L(c_i)$:linguistic signal

Why this formula?

- Semantic similarity alone misses zero-day attacks
- Linguistic signals alone are too weak
- Combining both ensures detection of **explicit + implicit instructions**

Module 6: Cross Chunk Collusion Analysis

Purpose:

- Detect **emergent attacks** formed by combinations of chunks.

Working


- Construct graph $G = (V, E)$
 - Nodes = chunks
 - Edges = semantic similarity between chunks
- Instruction signals propagate across connected chunks
- Detect:
 - Instruction chains
- Reinforcement patterns

$$Collusion(c_i) = \sum_{j \in Neighbors(i)} W_{ji} \cdot Instructionality(c_j)$$


Novelty:

- Existing RAG defenses **do not model chunk interactions**.

Cross-Chunk Interaction Graph
240 nodes | 3652 edges | Node size $\propto$ trust score | Red edges = forced collusion links



Out-Degree Distribution by Chunk Label
(Full Corpus Graph)



Module 7: Trust Weighted Context Injection

Purpose

- Assign reliability to each chunk before giving it to LLM

Working

- Combine:
 - Instructionality score
 - Collusion risk
 - Source provenance

Formula

$$Trust(c_i) = w_1 \cdot (1 - Instructionality(c_i)) + w_2 \cdot (1 - Collusion(c_i)) + w_3 \cdot Provenance(c_i)$$

Where

- w_1, w_2, w_3 : weights
- Provenance = source credibility

Module 8: Groundedness Deviation Detection

Purpose:

- Catch failures **after generation**.

Checks Performed:

- Are output claims supported by high trust chunks?
- Did excluded or demoted chunks influence output?

Mitigation:

- Regenerate with stricter constraints
- Refuse if violation persists


Why Needed:

- Zero day attacks can bypass pre generation defenses.

Output

- The proposed runtime defense successfully separated benign content from adversarial content and maintained **zero false positives** in the frozen benchmark.
- In the main benchmark, the **5-layer system** achieved **Precision = 1.000**, **Recall = 0.450**, **F1 = 0.620**, **FPR = 0.000**, and **JSR = 0.550**, showing strong safety control with no benign-query blocking.
- After validation-based threshold tuning, performance improved substantially at the operating point, reaching **Precision = 1.000**, **Recall = 0.950**, **F1 = 0.974**, and **FPR = 0.000**.
- The system remained lightweight, with average latency around **26.3 ms** and p95 latency of **30.8 ms**, which supports interactive deployment.
- Ablation results showed that the intermediate defense layers, especially boundary detection and collusion analysis, contributed the most to overall effectiveness.

F1 vs instruct_threshold
(collusion_thr = 0.35)



Conclusion

01

Robustness: Runtime defense is more effective than static fine-tuning for RAG-specific threats.

02

Model Agnostic: Works equally well with GPT-4 (closed) and Llama-3/Qwen (open) because it focuses on the **context** rather than the weights.

03

Contribution: Addresses the "Instructional Authority" gap—ensuring the LLM treats retrieved data as *information* rather than *commands*.

Future Enhancement

01

Obfuscation Defense:

Developing adaptive filters for attacks hidden in Base64 encoding or complex paraphrasing.

02

Multimodal RAG:

Extending the pipeline to analyze images and audio files in multimodal RAG pipelines.

03

Red Teaming

Automation: Using "Attacker Agents" to continuously probe the defense layers for weaknesses.

04

Agentic Security:

Applying behavior auditing to autonomous agents that use RAG to perform actions (e.g., sending emails).

Outcome

- The project demonstrated that **runtime defense for RAG systems is practical and effective** without modifying or retraining the base LLM.
- It introduced a modular pipeline that can detect **instructional risk, cross-chunk collusion, and groundedness deviation** at inference time.
- The final system provides a **model-independent security layer** that can be applied to both open-weight and API-based LLMs.
- The evaluation confirmed that the approach can preserve utility while improving safety, especially after threshold calibration and trust-based injection tuning.
- Overall, the work validates a **safe and extensible direction for securing LLM-based RAG pipelines** against indirect prompt injection and jailbreak-style attacks.

References

Based on the source materials provided, here is the complete list of papers formatted as a reference list:

- 1.Xu, Z., Liu, Y., Deng, G., Li, Y., & Picek, S. (2024).** A Comprehensive Study of Jailbreak Attack versus Defense for Large Language Models. In *Findings of the Association for Computational Linguistics: ACL 2024*.
- 2.Hung, K. H., Ko, C. Y., Rawat, A., Chung, I. H., Hsu, W. H., & Chen, P. Y. (2025).** Attention Tracker: Detecting Prompt Injection Attacks in LLMs. In *Findings of the Association for Computational Linguistics: NAACL 2025*, 2309–2322.
- 3.Shen, X., Chen, Z., Backes, M., Yun, S., & Zhang, Y. (2024).** "Do Anything Now": Characterizing and Evaluating In-The-Wild Jailbreak Prompts on Large Language Models. In *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security (CCS 2024)*, 1671–1685.
- 4.Yi, S., Liu, Y., Sun, Z., Cong, T., He, X., Song, J., Xu, K., & Li, Q. (2024).** Jailbreak Attacks and Defenses Against Large Language Models: A Survey. *arXiv preprint arXiv:2407.04295*.
- 5.Yao, H., Shi, H., Chen, Y., Jiang, Y., Wang, C., & Qin, Z. (2025).** ControlNet: A Firewall for RAG-based LLM System. *arXiv preprint arXiv:2504.09593*.
- 6.Ramakrishnan, B., & Balaji, A. (2025).** Securing AI Agents Against Prompt Injection Attacks: A Comprehensive Benchmark and Defense Framework. *arXiv preprint arXiv:2511.15759*.
- 7.Qi, X., Panda, A., Lyu, K., Ma, X., Roy, S., Beirami, A., Mittal, P., & Henderson, P. (2024).** Safety Alignment Should Be Made More Than Just a Few Tokens Deep. *arXiv preprint arXiv:2406.05946*.
- 8.Wu, T., & Zhang, H. (2026).** Chain-of-Detection enables robust and efficient jailbreak defense. *Neural Networks*, Vol. 196.
- 9.Mazeika, M., Phan, L., Yin, X., Zou, A., Wang, Z., Mu, N., Sakhaee, E., Li, N., Basart, S., Li, B., Forsyth, D., & Hendrycks, D. (2024).** HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal. *arXiv preprint arXiv:2402.04249*.
- 10.[Project Thesis]. (2026).** Securing RAG Systems Against Indirect Prompt Injection Attacks: A Modular, Model-Independent Runtime Security Middleware for Retrieval-Augmented Generation.
- 11.Yu, J., Lin, X., Yu, Z., & Xing, X. (2024).** LLM-Fuzzer: Scaling Assessment of Large Language Model Jailbreaks. In *33rd USENIX Security Symposium*, 4657–4674.
- 12.Zou, A., Wang, Z., Kolter, J. Z., & Fredrikson, M. (2023).** Universal and Transferable Adversarial Attacks on Aligned Language Models. *arXiv preprint arXiv:2307.15043*.



DETECTION AND MITIGATION OF JAILBREAK ATTACKS IN LLM'S

School of Computer Science and Engineering – Winter Semester 2025 - 26

Introduction

Retrieval-Augmented Generation improves factual grounding, but it also creates a new attack surface where malicious instructions can be hidden inside retrieved documents. This project proposes a black-box runtime defense that inspects retrieved chunks before they reach the LLM, without retraining the model or accessing its internal weights. The design is modular, model-agnostic, and focused on runtime protection for API-based and open-weight RAG systems.

Motivation

RAG systems are vulnerable to **hidden instructions in retrieved data**; Attacks can be **distributed across multiple chunks** (collusion); Static filters fail against **zero-day jailbreaks and obfuscation**; API-based LLMs prevent access to internal signals or retraining

Goal: Design a **practical, deployable runtime defense** that improves security while preserving usability and low latency.

Scope of the Project

Detect instruction vs factual content in retrieved chunks; Identify **cross-chunk collusion using graph-based analysis**; Assign **trust scores** using instructionality, collusion risk, and provenance; Apply **trust-weighted context injection** instead of binary filtering; Validate outputs using **groundedness checking**

Constraints:

No modification or fine-tuning of LLM; Text-based RAG pipelines only; Focus on runtime efficiency and low false positives

Methodology

The system is implemented as a **five-layer runtime defense pipeline**:


Instruction–Data Boundary Detection
Detects hidden instructions using semantic similarity and linguistic cues

Cross-Chunk Collusion Analysis
Models chunk interactions using a bounded graph to detect distributed attacks

Trust Scoring Engine
Combines instructionality, collusion risk, and provenance

Trust-Weighted Context Injection
Controls influence of each chunk instead of allow/block filtering

Groundedness & Output Validation
Ensures generated responses rely only on trusted information



These signals are combined with provenance into a **trust score**, which controls **trust-weighted context injection**


$$Trust(c_i) = \omega_1 \cdot (1 - I(c_i)) + \omega_2 \cdot (1 - C(c_i)) + \omega_3 \cdot P(c_i)$$

$I(c_i)$:instructionality
 $C(c_i)$:collusion risk
 $P(c_i)$:provenance score


Results

The evaluation used **5,000 benign documents, 500 poisoned chunks, 200 collusive sets, and 1,000 attack queries**.


After validation-based calibration, the selected thresholds were **instructionality = 0.20** and **collusion = 0.35**. At this tuned operating point, the system reached **Precision = 1.000, Recall = 0.950, F1 = 0.974**, and **FPR = 0.000**, showing a strong separation between benign and adversarial samples.



The runtime overhead remained low enough for interactive use, with **mean latency = 26.3 ms** and **p95 latency = 30.8 ms**. This supports the key claim that the defense can be deployed at inference time without materially slowing the system.



The ablation study showed that the largest gains came from the **boundary detection** and **collusion analysis** layers. Layer 1 alone performed weakly, but adding Layer 3 gave a major jump in performance, and Layer 4 improved detection further by capturing cross-chunk reinforcement. The final validation layer improved safety behavior, although it also made the overall benchmark more conservative.



Conclusion

Results show strong separation between benign and malicious inputs, achieving **high F1 score after calibration (0.974)** while maintaining **zero false positives** and low runtime overhead (~26 ms).

The results support the hypothesis that a **model-agnostic, runtime defense** can enhance RAG security without modifying the underlying LLM.

Improve robustness against paraphrased and obfuscated attacks.
Strengthen semantic detection and refine validation to reduce conservativeness.
Extend the framework to longer contexts and more complex multi-step interactions.

References

Zhang, P., et al. (2024). "Prompt Automatic Iterative Refinement (PAIR) for LLM Jailbreaking." *arXiv preprint*. (Automated Jailbreak Frameworks)
Rawat, A., et al. (2025). "Retrieval-Augmented Defense (RAD): Behavioral anomaly detection in RAG systems." *Security and Privacy in AI*. (RAG Security & Behavioral Firewalls)

Student RegNo, Name: 22BCE2036, Muskan Agarwal
22BCE3689, Divyesh Narshi Kanzariya
22BCI0286, Tharun Suresh Kumar
Guide Name: Prof. Umamaheswari M